


F7. Agentic RAG as controlled evidence acquisition

Policy learning changes when and how latent-chain evidence becomes observable.

Accuracy gain is unattributed unless
you report P/S/O/U/T diagnostics.



$$R_T = \mathbf{1}\{\hat{a} = a^*\} - \lambda_c \sum_t \text{cost}(u_t) - \lambda_f \mathbf{1}\{\hat{\mathcal{A}} \leftarrow C^*\}$$